

Machine Learning Course Project

Pedro Rojas G

Sep 26, 2020

Background

Using devices such as Jawbone Up, Nike FuelBand, and Fitbit it is now possible to collect a large amount of data about personal activity relatively inexpensively. These type of devices are part of the quantified self movement - a group of enthusiasts who take measurements about themselves regularly to improve their health, to find patterns in their behavior, or because they are tech geeks. One thing that people regularly do is quantify how much of a particular activity they do, but they rarely quantify how well they do it. In this project, your goal will be to use data from accelerometers on the belt, forearm, arm, and dumbbell of 6 participants. They were asked to perform barbell lifts correctly and incorrectly in 5 different ways

Task

The goal of this project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

1. Your submission should consist of a link to a Github repo with your R markdown and compiled HTML file describing your analysis. Please constrain the text of the writeup to < 2000 words and the number of figures to be less than 5. It will make it easier for the graders if you submit a repo with a gh-pages branch so the HTML page can be viewed online (and you always want to make it easy on graders).
2. You should also apply your machine learning algorithm to the 20 test cases available in the test data above. Please submit your predictions in appropriate format to the programming assignment for automated grading. See the programming assignment for additional details.

Data loading and cleaning

```
library(rattle)

## Rattle: A free graphical interface for data science with R.
## Version 5.2.0 Copyright (c) 2006-2018 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.

library(caret)

## Loading required package: lattice
```

```
## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 3.5.3
library(rpart)
library(rpart.plot)
library(corrplot)
## Warning: package 'corrplot' was built under R version 3.5.3
## corrplot 0.84 loaded
library(randomForest)
## randomForest 4.6-14
## Type rfNews() to see new features/changes/bug fixes.
##
## Attaching package: 'randomForest'
## The following object is masked from 'package:ggplot2':
##
##     margin
## The following object is masked from 'package:rattle':
##
##     importance
library(RColorBrewer)
```

Now this is test to see if the dataset is downloaded in folder and if not it download it.

```
trainUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
testUrl <- "https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"
trainFile <- "./data/pml-training.csv"
testFile <- "./data/pml-testing.csv"
if (!file.exists("./data")) {
  dir.create("./data")
}
if (!file.exists(trainFile)) {
  download.file(trainUrl, destfile = trainFile, method = "curl")
}
if (!file.exists(testFile)) {
  download.file(testUrl, destfile = testFile, method = "curl")
}
```

```

}
rm(trainUrl)
rm(testUrl)

```

and we now read the data files

```

trainRaw <- read.csv(trainFile)
testRaw <- read.csv(testFile)
dim(trainRaw)
## [1] 19622 160
dim(testRaw)
## [1] 20 160
rm(trainFile)
rm(testFile)

```

the data need some cleaning so we remove missing data points and ignore the variables that we are not interested as they are useless in the analysis.

```

NZV <- nearZeroVar(trainRaw, saveMetrics = TRUE)
head(NZV, 20)

```

##		freqRatio	percentUnique	zeroVar	nzv
## X		1.000000	100.00000000	FALSE	FALSE
## user_name		1.100679	0.03057792	FALSE	FALSE
## raw_timestamp_part_1		1.000000	4.26562022	FALSE	FALSE
## raw_timestamp_part_2		1.000000	85.53154622	FALSE	FALSE
## cvtd_timestamp		1.000668	0.10192641	FALSE	FALSE
## new_window		47.330049	0.01019264	FALSE	TRUE
## num_window		1.000000	4.37264295	FALSE	FALSE
## roll_belt		1.101904	6.77810621	FALSE	FALSE
## pitch_belt		1.036082	9.37722964	FALSE	FALSE
## yaw_belt		1.058480	9.97349913	FALSE	FALSE
## total_accel_belt		1.063160	0.14779329	FALSE	FALSE
## kurtosis_roll_belt		1921.600000	2.02323922	FALSE	TRUE
## kurtosis_picth_belt		600.500000	1.61553358	FALSE	TRUE
## kurtosis_yaw_belt		47.330049	0.01019264	FALSE	TRUE
## skewness_roll_belt		2135.111111	2.01304658	FALSE	TRUE
## skewness_roll_belt.1		600.500000	1.72255631	FALSE	TRUE

```
## skewness_yaw_belt      47.330049      0.01019264      FALSE      TRUE
## max_roll_belt          1.000000      0.99378249      FALSE      FALSE
## max_pitch_belt         1.538462      0.11211905      FALSE      FALSE
## max_yaw_belt           640.533333      0.34654979      FALSE      TRUE

training01 <- trainRaw[, !NZV$nzv]
testing01 <- testRaw[, !NZV$nzv]
dim(training01)
## [1] 19622 100
dim(testing01)
## [1] 20 100

rm(trainRaw)
rm(testRaw)
rm(NZV)
```

the variables we are going to remove is that the variables that does not contribute in accelerometer measurements.


```
regex <- grepl("^X|timestamp|user_name", names(training01))
training <- training01[, !regex]
testing <- testing01[, !regex]
rm(regex)
rm(training01)
rm(testing01)
dim(training)
## [1] 19622 95
dim(testing)
## [1] 20 95
```

and we finally remove NAN

```
cond <- (colSums(is.na(training)) == 0)
training <- training[, cond]
testing <- testing[, cond]
rm(cond)
```

now we visualize the correlation between different variables in the dataset. based on this we can see our ignorance of variables before if it is ok or not.

```
corrplot(cor(training[, -length(names(training))]), method = "color", tl.
cex = 0.5)
```



Approach

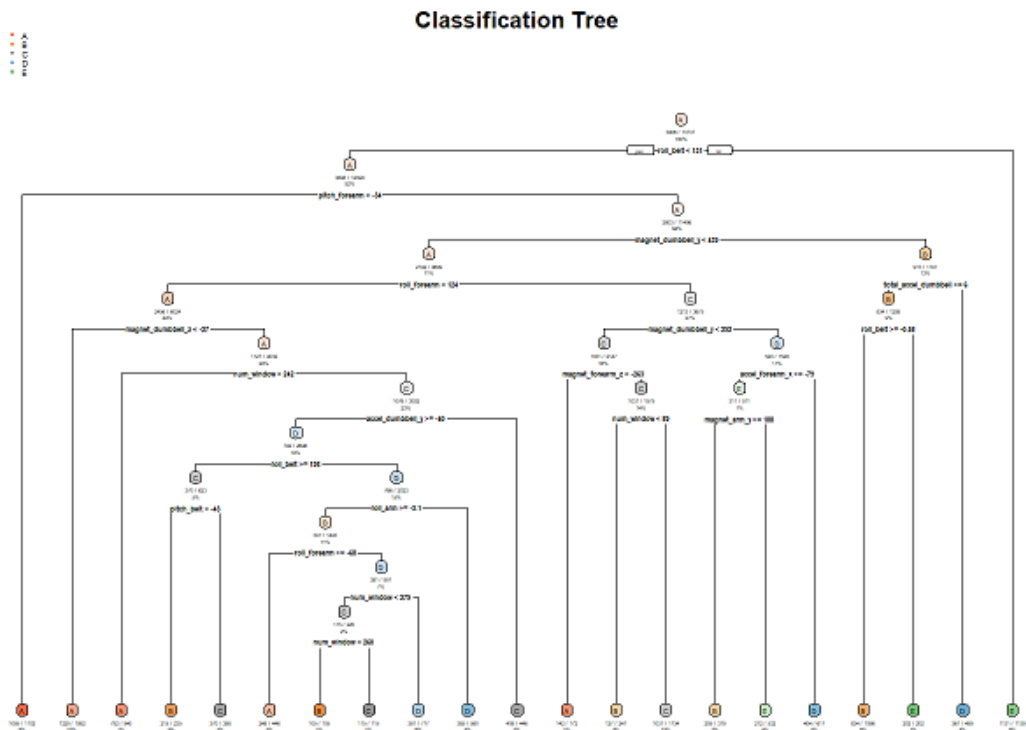
I am going to apply two different models and evaluate how they behave on this data. Two models will be run and they are decision tree and random forest. we seek the model with the highest accuracy will be our final model. we will use the ordinary way to split the cleaned training set into a pure training data set (70%) and a validation data set (30%). We will use the validation data set to conduct cross validation. We are using seed for reproducibility purposes.

```
set.seed(56789) # For reproducible purpose
inTrain <- createDataPartition(training$classe, p = 0.70, list = FALSE)
validation <- training[-inTrain, ]
training <- training[inTrain, ]
rm(inTrain)
```

Decision Tree

we are using decision tree to fit our model

```
modelTree <- rpart(classe ~ ., data = training, method = "class")
rpart.plot(modelTree, main="Classification Tree", extra=102, under=TRUE,
facilen=0)
```



Now after we have trained our model, we want to test it against validation data.

```
predictTree <- predict(modelTree, validation, type = "class")
confusionMatrix(validation$classe, predictTree)
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      A      B      C      D      E
```

```
##           A 1526    41    20    61    26
```

```
##           B  264   646    74   126    29
```

```
##           C   20    56   852    72    26
```

```
##           D    93    31   133   665    42
##           E    82    85    93   128   694
##
## Overall Statistics
##
##           Accuracy : 0.7448
##           95% CI : (0.7334, 0.7559)
##           No Information Rate : 0.3373
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.6754
##           McNemar's Test P-Value : < 2.2e-16
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.7688    0.7520    0.7270    0.6321    0.8494
## Specificity           0.9621    0.9019    0.9631    0.9381    0.9234
## Pos Pred Value        0.9116    0.5672    0.8304    0.6898    0.6414
## Neg Pred Value        0.8910    0.9551    0.9341    0.9214    0.9744
## Prevalence            0.3373    0.1460    0.1992    0.1788    0.1388
## Detection Rate        0.2593    0.1098    0.1448    0.1130    0.1179
## Detection Prevalence  0.2845    0.1935    0.1743    0.1638    0.1839
## Balanced Accuracy     0.8654    0.8270    0.8450    0.7851    0.8864
accuracy <- postResample(predictTree, validation$classe)
ose <- 1 - as.numeric(confusionMatrix(validation$classe, predictTree)$overall[1])
rm(predictTree)
rm(modelTree)
```

We find that the Estimated Accuracy of the Decision tree Model is 74.4774851% and the Estimated Out-of-Sample Error is about 25.5225149%.

Random forest

We now train our model using random forest and doing the same validation

```

modelRF <- train(classe ~ ., data = training, method = "rf", trControl =
trainControl(method = "cv", 5), ntree = 250)

modelRF

## Random Forest
##
## 13737 samples
##    53 predictor
##    5 classes: 'A', 'B', 'C', 'D', 'E'
##
## No pre-processing
## Resampling: Cross-Validated (5 fold)
## Summary of sample sizes: 10989, 10990, 10990, 10990, 10989
## Resampling results across tuning parameters:
##
##   mtry  Accuracy   Kappa
##   2     0.9949768  0.9936459
##   27    0.9976705  0.9970535
##   53    0.9957051  0.9945672
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 27.

```

Now after we have trained our model, we want to test it against validation data.

```

predictRF <- predict(modelRF, validation)
confusionMatrix(validation$classe, predictRF)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      A      B      C      D      E
##      A 1674      0      0      0      0
##      B   3 1136      0      0      0
##      C   0   1 1022      3      0
##      D   0   0   4  960      0
##      E   0   0   0   1 1081
##

```



```
## Overall Statistics
##
##           Accuracy : 0.998
##           95% CI : (0.9964, 0.9989)
##           No Information Rate : 0.285
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9974
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity           0.9982   0.9991   0.9961   0.9959   1.0000
## Specificity           1.0000   0.9994   0.9992   0.9992   0.9998
## Pos Pred Value        1.0000   0.9974   0.9961   0.9959   0.9991
## Neg Pred Value         0.9993   0.9998   0.9992   0.9992   1.0000
## Prevalence            0.2850   0.1932   0.1743   0.1638   0.1837
## Detection Rate         0.2845   0.1930   0.1737   0.1631   0.1837
## Detection Prevalence   0.2845   0.1935   0.1743   0.1638   0.1839
## Balanced Accuracy      0.9991   0.9992   0.9976   0.9975   0.9999
accuracy <- postResample(predictRF, validation$classe)
ose <- 1 - as.numeric(confusionMatrix(validation$classe, predictRF)$overall[1])
rm(predictRF)
```

We find that the Estimated Accuracy of the Random Forest Model is 99.7960918% and the Estimated Out-of-Sample Error is about 0.2039082%.

Conlusion

we find that the Accuracy of the Random Forest Model and error is better than the Decision Tree model. so we conclude that the random forest is the better model.

submission part

this is the code for predicting outcome levels on the original Testing data set using Random Forest algorithm as it is the chosn model as being better at performance on our data.

```
rm(accuracy)
rm(ose)
predict(modelRF, testing[, -length(names(testing))])
##  [1] B A B A A E D B A A B C B A E E A B B B
## Levels: A B C D E
```